

# PART A: TECHNICAL DESIGN DOCUMENT

## ML Production System for App Similarity & Performance Prediction

MobUpps Home Assignment

Author: Ortal Lasry | ML Engineer

Date: October 2025

Version: 1.0

### 1. SYSTEM ARCHITECTURE OVERVIEW

#### 1.1 High-Level Architecture

The system is designed as a microservices architecture on AWS, capable of scaling from 100 to 1M requests/day with minimal latency (<200ms p99).

#### Key Components:

- API Gateway: Entry point for all client requests
- ECS Fargate: Containerized application hosting (auto-scaling)
- ElastiCache Redis: Embedding cache + session storage
- S3: Data lake for embeddings, metadata, and performance logs
- DynamoDB: Real-time metrics and A/B test assignments
- CloudWatch: Monitoring, logging, and alerting
- SageMaker: Model retraining and versioning
- Lambda: Event-driven processing (feedback loop)
- Step Functions: Orchestration for model deployment pipeline

#### 1.2 Data Flow

##### [Client Request]

- > API Gateway (Rate limiting, auth)
- > Application Load Balancer
- > ECS Fargate (Similarity Service)
- > ElastiCache (Check embedding cache)
- > If miss: S3 (Load embeddings)

```
-> DynamoDB (Log A/B assignment + metrics)
-> Return similar apps + predictions
<- Response to Client
```

### [Feedback Loop]

```
-> CloudWatch Logs (Capture predictions + outcomes)
-> Kinesis Data Streams (Real-time event processing)
-> Lambda (Aggregate performance metrics)
-> S3 (Store training data)
-> SageMaker (Model retraining trigger)
-> Model Registry (Version new embeddings)
-> Step Functions (Orchestrate deployment)
```

## 1.3 Architecture Diagram Description

### Layer 1 - Client Layer:

- Mobile Apps / Web Clients
- SDK/API Integration

### Layer 2 - Edge Layer (AWS CloudFront + API Gateway):

- CloudFront: CDN for static assets + edge caching
- API Gateway: REST API with request validation
- WAF: DDoS protection + rate limiting
- Cognito: Authentication (if needed)

### Layer 3 - Application Layer (ECS Fargate):

- Cluster: Auto-scaling (2-50 tasks)
- Service 1: Similarity API (FastAPI containers)
- Service 2: Prediction API (FastAPI containers)
- Load Balancer: ALB with health checks
- Target Groups: Blue/Green deployment support

### Layer 4 - Data Layer:

- ElastiCache Redis:
- \* Embedding cache (TTL: 1 hour)
- \* A/B test assignments (sticky sessions)
- \* Rate limiting counters

- DynamoDB:
  - \* Metrics table (real-time aggregations)
  - \* A/B test results (partition by experiment\_id)
  - \* Request logs (time-series data)
- S3:
  - \* Embeddings bucket (v1, v2, v3...)
  - \* Historical data bucket (training datasets)
  - \* Logs bucket (CloudWatch exports)

## Layer 5 - ML Pipeline (SageMaker + Step Functions):

- SageMaker Training: Model retraining jobs
- SageMaker Endpoint: Real-time inference (if needed)
- Model Registry: Versioned embeddings
- Step Functions: Deploy workflow automation

## Layer 6 - Monitoring & Observability:

- CloudWatch Logs: Structured logging
- CloudWatch Metrics: Custom metrics (latency, accuracy)
- CloudWatch Alarms: Threshold-based alerts
- X-Ray: Distributed tracing
- SNS: Alert notifications (email, Slack, PagerDuty)

## 2. AWS SERVICES SELECTION & JUSTIFICATION

### 2.1 Core Services

- Amazon ECS Fargate\*\* (Application Hosting)

#### Why:

- Serverless containers (no EC2 management)
- Auto-scaling based on CPU/memory/custom metrics
- Pay per vCPU-second (cost-efficient for variable load)
- Blue/Green deployments with ALB
- Integrates with CloudWatch for observability

Alternative Considered: EKS (more complex), Lambda (cold start issues)

- Amazon API Gateway\*\* (API Management)

### **Why:**

- Built-in rate limiting (10,000 req/sec burst)
- Request/response validation
- API versioning support
- AWS WAF integration for security
- Usage plans for different client tiers

Alternative Considered: ALB only (lacks API management features)

- Amazon ElastiCache Redis\*\* (Caching + Session Storage)

### **Why:**

- Sub-millisecond latency for embedding lookups
- Cluster mode for horizontal scaling
- Automatic failover (Multi-AZ)
- Supports complex data structures (hashes for embeddings)
- TTL for cache invalidation

Alternative Considered: DynamoDB DAX (higher cost), Memcached (no persistence)

- Amazon DynamoDB\*\* (Real-time Metrics + A/B Assignments)

### **Why:**

- Single-digit millisecond latency at scale
- Auto-scaling (pay per request)
- Global Secondary Indexes for querying by experiment\_id
- TTL for automatic data cleanup
- DynamoDB Streams for change data capture

Alternative Considered: RDS Aurora (more expensive, over-engineered)

- Amazon S3\*\* (Data Lake)

### **Why:**

- Unlimited storage capacity
- 99.999999999% durability
- Lifecycle policies (move to Glacier after 90 days)
- S3 Select for querying CSV/Parquet
- Event notifications (trigger Lambda on new embeddings)

Alternative Considered: EFS (higher cost for large datasets)

- Amazon SageMaker\*\* (Model Management)

### **Why:**

- Model Registry for versioning
- Training jobs with spot instances (70% cost savings)
- Batch transform for bulk predictions
- Feature Store for metadata management
- Pipelines for end-to-end ML workflows

Alternative Considered: Self-managed MLflow (operational overhead)

- AWS Step Functions\*\* (Orchestration)

### **Why:**

- Visual workflow designer
- Error handling and retries
- Integration with all AWS services
- Audit trail for deployments
- State machine versioning

Alternative Considered: Airflow on EC2 (requires maintenance)

## 2.2 Monitoring & Observability

- Amazon CloudWatch\*\* (Logging + Metrics + Alarms)

### **Why:**

- Native AWS integration (no agents needed)
- Log Insights for querying structured logs
- Custom metrics (business KPIs)
- Anomaly detection using ML
- Dashboard for real-time monitoring

Cost: ~\$0.50/GB ingested + \$0.03/GB analyzed

- AWS X-Ray\*\* (Distributed Tracing)

### **Why:**

- Trace requests across microservices
- Identify bottlenecks (slow S3 calls, Redis misses)
- Service map visualization
- Integration with CloudWatch ServiceLens

Cost: ~\$5 per million traces

- Amazon SNS\*\* (Alerting)

### **Why:**

- Fan-out to multiple channels (email, Slack, PagerDuty)
- Message filtering
- Mobile push notifications
- Dead-letter queue for failed deliveries

Cost: ~\$0.50 per million notifications

## 3. A/B TESTING STRATEGY

### 3.1 Traffic Routing Strategy

- Approach: Hash-Based Sticky Sessions\*\*

### **Implementation:**

1. Hash partner\_id + app\_id using MD5
2. Map hash to [0, 100] range
3. Route based on configured split (e.g., 50/50)
4. Store assignment in Redis (TTL: 24 hours)

### **Benefits:**

- Deterministic: Same user always gets same variant
- Stateless: No database lookup for every request
- Configurable: Adjust split via environment variables
- Fast: O(1) assignment decision

### **Code Example:**

#### **def pick\_arm(partner\_id, app\_id, v1\_weight=0.5):**

```
key = f"{partner_id}:{app_id}"
hash_val = int(hashlib.md5(key.encode()).hexdigest(), 16) % 100
return "v1" if hash_val < v1_weight * 100 else "v2"
```

### 3.2 A/B Test Configuration

DynamoDB Table: ab\_experiments

- experiment\_id (PK): "embedding\_v1\_vs\_v2\_2025\_10"

- status: "active" | "completed" | "stopped"
- v1\_weight: 0.5 (50% traffic)
- v2\_weight: 0.5 (50% traffic)
- start\_date: "2025-10-15"
- end\_date: "2025-10-30"
- winning\_criteria: "ctr\_improvement > 5%"

DynamoDB Table: ab\_metrics

- experiment\_id (PK): "embedding\_v1\_vs\_v2\_2025\_10"
- timestamp (SK): 1729360000
- arm: "v1" | "v2"
- requests\_count: 10000
- avg\_latency\_ms: 25.3
- p99\_latency\_ms: 89.1
- ctr: 0.0235
- prediction\_error: 0.0123

### 3.3 Metrics to Track

- Performance Metrics:\*\*
  - Latency (p50, p95, p99)
  - Throughput (requests/second)
  - Error rate (4xx, 5xx)
  - Cache hit rate
- Business Metrics:\*\*
  - Click-Through Rate (CTR)
  - Conversion Rate
  - Prediction accuracy (MAE, RMSE)
  - User engagement
- Statistical Significance:\*\*
  - Use Bayesian A/B testing (Beta distribution)
  - Require minimum sample size: 10,000 requests per arm
  - Confidence level: 95%
  - Minimum detectable effect: 5% improvement

### 3.4 Winner Selection Logic

## Automated Decision (Lambda triggered daily):

1. Check if experiment has run for minimum duration (7 days)
2. Verify minimum sample size reached (10K per arm)
3. Calculate statistical significance (Bayesian credible interval)

### 4. If winner detected:

- `ctr_improvement > 5% AND p_value < 0.05`
- Gradually shift traffic to winner (10% increments)
- Monitor for regressions
- After 3 days of stability, promote to 100%

5. Notify team via SNS (Slack channel)

## 4. CI/CD PIPELINE OVERVIEW

### 4.1 Pipeline Stages

#### Stage 1: Source (GitHub)

- Push to main branch triggers pipeline
- GitHub Actions webhook to AWS CodePipeline
- Or use AWS CodeCommit for native integration

#### Stage 2: Build (AWS CodeBuild)

- Pull Docker base image from ECR
- Run unit tests (pytest)
- Run integration tests (testcontainers)
- Run performance tests (k6 or locust)
- Build Docker image
- Tag with commit SHA + timestamp
- Push to Amazon ECR

#### Stage 3: Deploy to Staging (ECS Fargate)

- Update ECS task definition (new image tag)
- Deploy to staging cluster
- Run smoke tests
- Run end-to-end tests
- Performance benchmarking

#### Stage 4: Manual Approval (Optional)



- Notify team via SNS
- Review test results
- Approve/Reject via AWS Console or CLI

#### Stage 5: Deploy to Production (Blue/Green)

- Create new ECS task set (Green)
- Route 10% traffic to Green (canary)
- Monitor metrics for 10 minutes
- If healthy: shift remaining 90%
- If unhealthy: automatic rollback to Blue
- Terminate old task set after 1 hour

#### Stage 6: Post-Deployment Validation

- Run synthetic tests (CloudWatch Synthetics)
- Monitor error rate, latency
- Alert if anomalies detected

### 4.2 CodePipeline Configuration

#### **buildspec.yml (CodeBuild):**

```
version: 0.2 phases: pre_build: commands: - echo Logging in to Amazon ECR... - aws ecr
get-login-password | docker login --username AWS ... - pip install -r requirements.txt build:
commands: - echo Running tests... - pytest tests/ --cov=app --cov-report=xml - echo Building
Docker image... - docker build -t similarity-service:$CODEBUILD_RESOLVED_SOURCE_VERSION . -
docker tag similarity-service:$CODEBUILD_RESOLVED_SOURCE_VERSION \
$ECR_REPO_URI:$CODEBUILD_RESOLVED_SOURCE_VERSION post_build: commands: - echo Pushing Docker
image... - docker push $ECR_REPO_URI:$CODEBUILD_RESOLVED_SOURCE_VERSION - echo Writing image
definitions file... - printf '["name":"similarity-api","imageUri":"%s"]' \
$ECR_REPO_URI:$CODEBUILD_RESOLVED_SOURCE_VERSION > imagedefinitions.json artifacts: files:
imagedefinitions.json
```

### 4.3 Deployment Strategies

- Blue/Green Deployment\*\* (Production)
- Two environments: Blue (current), Green (new)
- Instant rollback capability
- Zero downtime
- Cost: 2x resources during deployment (~10 minutes)
- Canary Deployment\*\* (High-Risk Changes)
- Route 5% -> 10% -> 25% -> 50% -> 100%
- Automated rollback if metrics degrade
- Slower but safer

- Rolling Update\*\* (Low-Risk Changes)
- Update tasks one-by-one
- No extra resources needed
- Longer deployment time

## 5. MONITORING & ALERTING

### 5.1 Metrics to Monitor

- Application Metrics (CloudWatch Custom Metrics):\*\*
  - similarity\_search\_latency (ms) [p50, p95, p99]
  - prediction\_latency (ms)
  - embedding\_cache\_hit\_rate (%)
  - ab\_test\_assignment\_latency (ms)
  - requests\_per\_second (count)
  - error\_rate\_4xx (%)
  - error\_rate\_5xx (%)
- Infrastructure Metrics (CloudWatch Default):\*\*
  - ECS CPU Utilization (%)
  - ECS Memory Utilization (%)
  - ALB Target Response Time (seconds)
  - ALB Healthy Host Count
  - Redis CPU Utilization (%)
  - Redis Evictions (count)
  - DynamoDB Read/Write Capacity (units)
  - API Gateway 4XXError, 5XXError (count)
- Business Metrics (Custom):\*\*
  - ctr\_by\_arm (v1, v2)
  - prediction\_accuracy\_mae
  - revenue\_per\_request (\$)
  - user\_satisfaction\_score

### 5.2 Alerting Strategy

- Critical Alerts (PagerDuty 24/7):\*\*
  - Error rate > 5% for 5 minutes

- p99 latency > 500ms for 10 minutes
- All ECS tasks unhealthy
- DynamoDB throttling (>100 events/min)
- Zero traffic for 5 minutes

- Warning Alerts (Slack #monitoring):\*\*
- Error rate > 2% for 10 minutes
- p99 latency > 300ms for 15 minutes
- Cache hit rate < 80%
- Redis memory > 80%
- Cost anomaly detected (>20% spike)

- Info Alerts (Email):\*\*
- Daily metrics summary
- Weekly A/B test report
- Monthly cost report

### 5.3 CloudWatch Alarms Configuration

#### Alarm: HighErrorRate

Metric: error\_rate\_5xx

Threshold: > 5%

Period: 5 minutes

Evaluation Periods: 2

Action: SNS -> PagerDuty

Auto-Scaling Action: Scale out by 2 tasks

#### Alarm: HighLatency

Metric: similarity\_search\_latency\_p99

Threshold: > 500ms

Period: 10 minutes

Evaluation Periods: 2

Action: SNS -> Slack

Auto-Scaling Action: Scale out by 1 task

#### Alarm: LowCacheHitRate

Metric: embedding\_cache\_hit\_rate

Threshold: < 80%

Period: 15 minutes

Evaluation Periods: 3

Action: SNS -> Email (investigate cache config)

## 5.4 Dashboard Layout (CloudWatch)

### Dashboard: Similarity Service Production

#### Panel 1: Traffic

- Requests/second (line chart)
- Error rate (stacked area: 4xx, 5xx)
- Cache hit rate (line chart)

#### Panel 2: Latency

- p50, p95, p99 latency (line chart)
- Latency heatmap (hour x percentile)

#### Panel 3: Infrastructure

- ECS CPU/Memory utilization (gauge)
- Healthy host count (number)
- Redis connections (line chart)

#### Panel 4: A/B Testing

- Traffic split (pie chart: v1 vs v2)
- CTR by arm (bar chart)
- Prediction accuracy by arm (line chart)

#### Panel 5: Business Metrics

- Revenue/hour (line chart)
- Cost/1M requests (number)

## 6. SCALING CONSIDERATIONS

### 6.1 Current State (100 requests/day)

#### Architecture:

- ECS Fargate: 1 task (0.25 vCPU, 0.5 GB RAM)
- ElastiCache: cache.t3.micro (0.5 GB)
- DynamoDB: On-demand pricing

- API Gateway: Pay-per-request
- S3: Standard storage

Cost Estimate: ~\$50/month

### **Performance:**

- Avg latency: ~30ms
- p99 latency: ~100ms
- Throughput: 100 req/day (~0.001 req/sec)

6.2 Medium Scale (10,000 requests/day)

### **Architecture:**

- ECS Fargate: 2-4 tasks (auto-scaling)
- ElastiCache: cache.t3.small (1.5 GB)
- DynamoDB: On-demand (auto-scaling)
- API Gateway: Pay-per-request
- S3: Standard storage

### **Auto-Scaling Policy:**

- Target CPU: 70%
- Target Memory: 80%
- Scale-out cooldown: 60 seconds
- Scale-in cooldown: 300 seconds

Cost Estimate: ~\$200/month

### **Performance:**

- Avg latency: ~30ms
- p99 latency: ~150ms
- Throughput: 10K req/day (~0.12 req/sec)

6.3 High Scale (1M requests/day)

### **Architecture:**

- ECS Fargate: 10-50 tasks (auto-scaling + scheduled scaling)
- ElastiCache: cache.r6g.xlarge (26 GB) - Cluster mode

- DynamoDB: Provisioned capacity with auto-scaling
- API Gateway: Regional with CloudFront CDN
- S3: Intelligent-Tiering

### **Optimizations:**

1. CloudFront edge caching for repeated queries
2. Redis cluster mode (3 shards) for horizontal scaling
3. DynamoDB Global Tables (multi-region if needed)
4. Batch processing for non-realtime predictions
5. Connection pooling (max 1000 connections per task)

### **Auto-Scaling Policy:**

- Target CPU: 60%
- Target Custom Metric: requests\_per\_task < 100/sec
- Scheduled Scaling: Scale up 2 hours before peak traffic
- Minimum tasks: 10 (always warm)
- Maximum tasks: 50 (cost limit)

Cost Estimate: ~\$2,000-3,000/month

### **Performance:**

- Avg latency: ~30ms
- p99 latency: ~200ms
- Throughput: 1M req/day (~12 req/sec avg, 50 req/sec peak)

6.4 Extreme Scale (10M+ requests/day)

### **Architecture Changes:**

- Multi-region deployment (us-east-1, eu-west-1)
- Global Accelerator for optimal routing
- Aurora Global Database for cross-region replication
- Lambda@Edge for edge computing
- Kinesis for event streaming (instead of direct DynamoDB writes)
- SQS for asynchronous predictions

Cost Estimate: ~\$10,000-15,000/month

## **Performance:**

- Avg latency: ~30ms
- p99 latency: ~200ms
- Throughput: 10M req/day (~120 req/sec avg, 500 req/sec peak)

### 6.5 Bottleneck Analysis & Solutions

#### Bottleneck 1: Embedding Cache Misses

Problem: Loading embeddings from S3 takes 100-500ms

#### **Solution:**

- Pre-warm cache on deployment
- Increase cache size (cache.r6g.4xlarge - 104 GB)
- Use Redis read replicas for read scaling

#### Bottleneck 2: Cosine Similarity Computation

Problem: Brute-force similarity search doesn't scale to 100K+ apps

#### **Solution:**

- Migrate to FAISS or Annoy for ANN search
- Use SageMaker endpoint for GPU-accelerated search
- Pre-compute top-1000 neighbors for popular apps

#### Bottleneck 3: DynamoDB Throttling

Problem: Burst traffic exceeds provisioned capacity

#### **Solution:**

- Use on-demand pricing (auto-scales instantly)
- Or: Set auto-scaling target to 70% utilization
- Use DynamoDB Streams to offload aggregations

#### Bottleneck 4: Cold Start Latency

Problem: New ECS tasks take 30-60s to start

#### **Solution:**

- Keep minimum 5 tasks always running
- Use Fargate Spot for cost savings (70% discount)
- Implement health check warm-up period

## 7. COST OPTIMIZATION STRATEGIES

## 7.1 Compute Optimization

- ECS Fargate Savings:\*\*
- Use Fargate Spot for non-critical tasks (70% savings)
- Right-size tasks (start with 0.25 vCPU, monitor utilization)
- Use Graviton2 (ARM) for 20% cost savings
- Set aggressive scale-in policy (shutdown idle tasks faster)

Estimated Savings: 40-50%

- Lambda for Batch Processing:\*\*
- Use Lambda for infrequent tasks (model evaluation, metrics aggregation)
- 1M free requests/month
- Pay per 100ms execution time

Estimated Savings: \$200/month vs. dedicated ECS task

## 7.2 Storage Optimization

- S3 Lifecycle Policies:\*\*
- Move embeddings to S3 Glacier after 90 days (80% savings)
- Move logs to S3 Glacier Deep Archive after 180 days (90% savings)
- Delete old training data after 1 year
- Use S3 Intelligent-Tiering for unpredictable access patterns

Estimated Savings: 60% on storage costs

- ElastiCache Reserved Instances:\*\*
- Purchase 1-year reserved instances for baseline capacity (35% savings)
- Use on-demand for burst capacity

Estimated Savings: 30-35% on cache costs

## 7.3 Data Transfer Optimization

- Reduce Cross-AZ Transfer:\*\*
- Use single-AZ deployment for dev/staging
- Use VPC endpoints for S3/DynamoDB (no internet gateway fees)
- Enable CloudFront compression (reduce bandwidth by 50%)



Estimated Savings: \$100-500/month at scale

#### 7.4 Monitoring Optimization

- CloudWatch Logs Optimization:\*\*
- Filter logs before ingestion (reduce noise)
- Use log sampling for high-volume logs (1% sample)
- Export to S3 after 7 days (90% savings)
- Use CloudWatch Logs Insights instead of Athena queries

Estimated Savings: 70% on log costs

#### 7.5 Cost Monitoring

- AWS Cost Explorer:\*\*
- Set budgets for each service
- Enable anomaly detection (detect cost spikes)
- Tag resources by environment (prod, staging, dev)
  
- Cost Allocation Tags:\*\*
- Team: data-science, ml-engineering
- Environment: prod, staging, dev
- Project: similarity-service
- CostCenter: 12345
  
- Monthly Cost Review:\*\*
- Identify top 10 cost drivers
- Right-size over-provisioned resources
- Delete unused resources (old ECS task definitions, AMLs)

#### 7.6 Total Cost Breakdown (1M requests/day)

ECS Fargate (20 tasks avg): \$800/month

ElastiCache (r6g.xlarge): \$400/month

DynamoDB (provisioned): \$300/month

S3 (1 TB storage): \$23/month

API Gateway (1M requests): \$3.50/month

CloudWatch (logs + metrics): \$200/month

Data Transfer: \$100/month

Load Balancer: \$20/month

Route53: \$1/month

SNS: \$1/month

Total: \$1,848/month

• With Optimizations:\*\*

• Fargate Spot: \$800 -> \$400 (-50%)

• ElastiCache RI: \$400 -> \$260 (-35%)

• S3 Lifecycle: \$23 -> \$10 (-57%)

• CloudWatch sampling: \$200 -> \$60 (-70%)

Optimized Total: \$1,055/month

• Annual Savings: \$9,516 (43% reduction)\*\*

## 8. FEEDBACK LOOP FOR MODEL IMPROVEMENT

### 8.1 Data Collection

#### **Real-Time Events (Kinesis Data Streams):**

- Prediction event: {app\_id, query, neighbors, predicted\_score, timestamp}
- Click event: {app\_id, campaign\_id, clicked, timestamp}
- Conversion event: {app\_id, campaign\_id, converted, revenue, timestamp}

#### **Batch Exports (Daily):**

- Export CloudWatch Logs to S3 (parquet format)
- Export DynamoDB metrics to S3 (time-series data)

### 8.2 Feature Store (SageMaker Feature Store)

Feature Group: app\_features

- app\_id (PK)
- category
- region
- pricing\_model
- features\_list
- historical\_ctr (updated daily)
- avg\_rating (updated weekly)
- last\_updated (timestamp)

## 8.3 Model Retraining Pipeline

Trigger: Weekly (or on-demand)

### Step 1: Data Preparation (Lambda)

- Query last 30 days of prediction + outcome data
- Join with feature store
- Filter outliers
- Save to S3 training bucket

### Step 2: Model Training (SageMaker Training Job)

- Load training data from S3
- Train new embedding model (v3)
- Validate on holdout set (20%)
- Compute metrics: recall@10, precision@10, MAP
- Save model artifacts to S3

### Step 3: Model Evaluation (Lambda)

- Compare v3 vs v2 on validation set
- If v3 improves recall@10 by >3%:
  - Promote to staging
  - Notify team via SNS
- Else:
  - Archive model
  - Log metrics to DynamoDB

### Step 4: Staging Deployment (Step Functions)

- Deploy v3 to staging environment
- Run regression tests
- Run A/B test simulator (synthetic traffic)
- If all tests pass:
  - Approve for production
- Else:
  - Rollback

### Step 5: Production Deployment (Gradual)

- Create new A/B experiment: v2 vs v3
- Route 10% traffic to v3

- Monitor for 48 hours
- If CTR improvement >5%:
  - Gradually shift to 50/50
  - Eventually 100% v3
- Deprecate v2 after 30 days

## 8.4 Monitoring Drift

### Feature Drift Detection:

- Monitor distribution of input features (category, region)
- Alert if >20% shift in 7 days
- Use SageMaker Model Monitor

### Prediction Drift Detection:

- Compare predicted\_ctr vs actual\_ctr
- Alert if MAE increases by >50%
- Trigger model retraining

## 9. SECURITY & COMPLIANCE

### 9.1 Security Best Practices

- Network Security:\*\*
  - VPC with private subnets for ECS tasks
  - Security groups: Allow only ALB -> ECS on port 8000
  - No public IPs for ECS tasks
  - VPC endpoints for AWS services (no internet gateway)
  - AWS WAF rules: Rate limiting, IP blocking, SQL injection protection
- Data Security:\*\*
  - S3 buckets: Block public access
  - S3 server-side encryption (SSE-S3 or KMS)
  - DynamoDB encryption at rest (KMS)
  - ElastiCache encryption in transit (TLS)
  - Secrets Manager for API keys, DB passwords
- Access Control:\*\*

- IAM roles for ECS tasks (least privilege)
- IAM roles for Lambda (separate per function)
- MFA for AWS Console access
- CloudTrail for audit logs
- Compliance:\*\*
- GDPR: PII data anonymization, data retention policies
- HIPAA: If handling health data, use HIPAA-eligible services
- SOC 2: Implement logging, monitoring, access controls

## 9.2 Disaster Recovery

- Backup Strategy:\*\*
- S3: Versioning enabled + Cross-Region Replication
- DynamoDB: Point-in-time recovery (PITR) + On-demand backups
- ElastiCache: Daily snapshots (7-day retention)
- RTO/RPO:\*\*
- RTO (Recovery Time Objective): < 1 hour
- RPO (Recovery Point Objective): < 5 minutes
- Multi-AZ deployment for high availability
- Failover Plan:\*\*
- 1. Detect outage (CloudWatch alarms)
- 2. Notify on-call engineer (PagerDuty)
- 3. Failover to secondary region (Route53 health check)
- 4. Investigate root cause
- 5. Post-mortem document

## 10. CONCLUSION

### **This architecture is designed to be:**

- Scalable: From 100 to 1M+ requests/day with auto-scaling
- Cost-Effective: \$50/month at low scale, \$1,000-3,000/month at high scale
- Reliable: 99.9% uptime with multi-AZ deployment
- Observable: Comprehensive monitoring and alerting
- Maintainable: CI/CD pipeline with automated testing
- Secure: VPC isolation, encryption at rest/transit, IAM least privilege

- ML-Ready: Feedback loop for continuous model improvement

The system leverages AWS managed services to minimize operational overhead and maximize developer productivity. The A/B testing framework enables data-driven decision making for model improvements.

### **Next Steps:**

1. Implement proof-of-concept on AWS Free Tier
2. Load test with k6 (simulate 1M requests/day)
3. Optimize based on bottleneck analysis
4. Document runbooks for on-call engineers
5. Train team on AWS services and monitoring dashboards

END OF DOCUMENT